

# Neural Networks 2

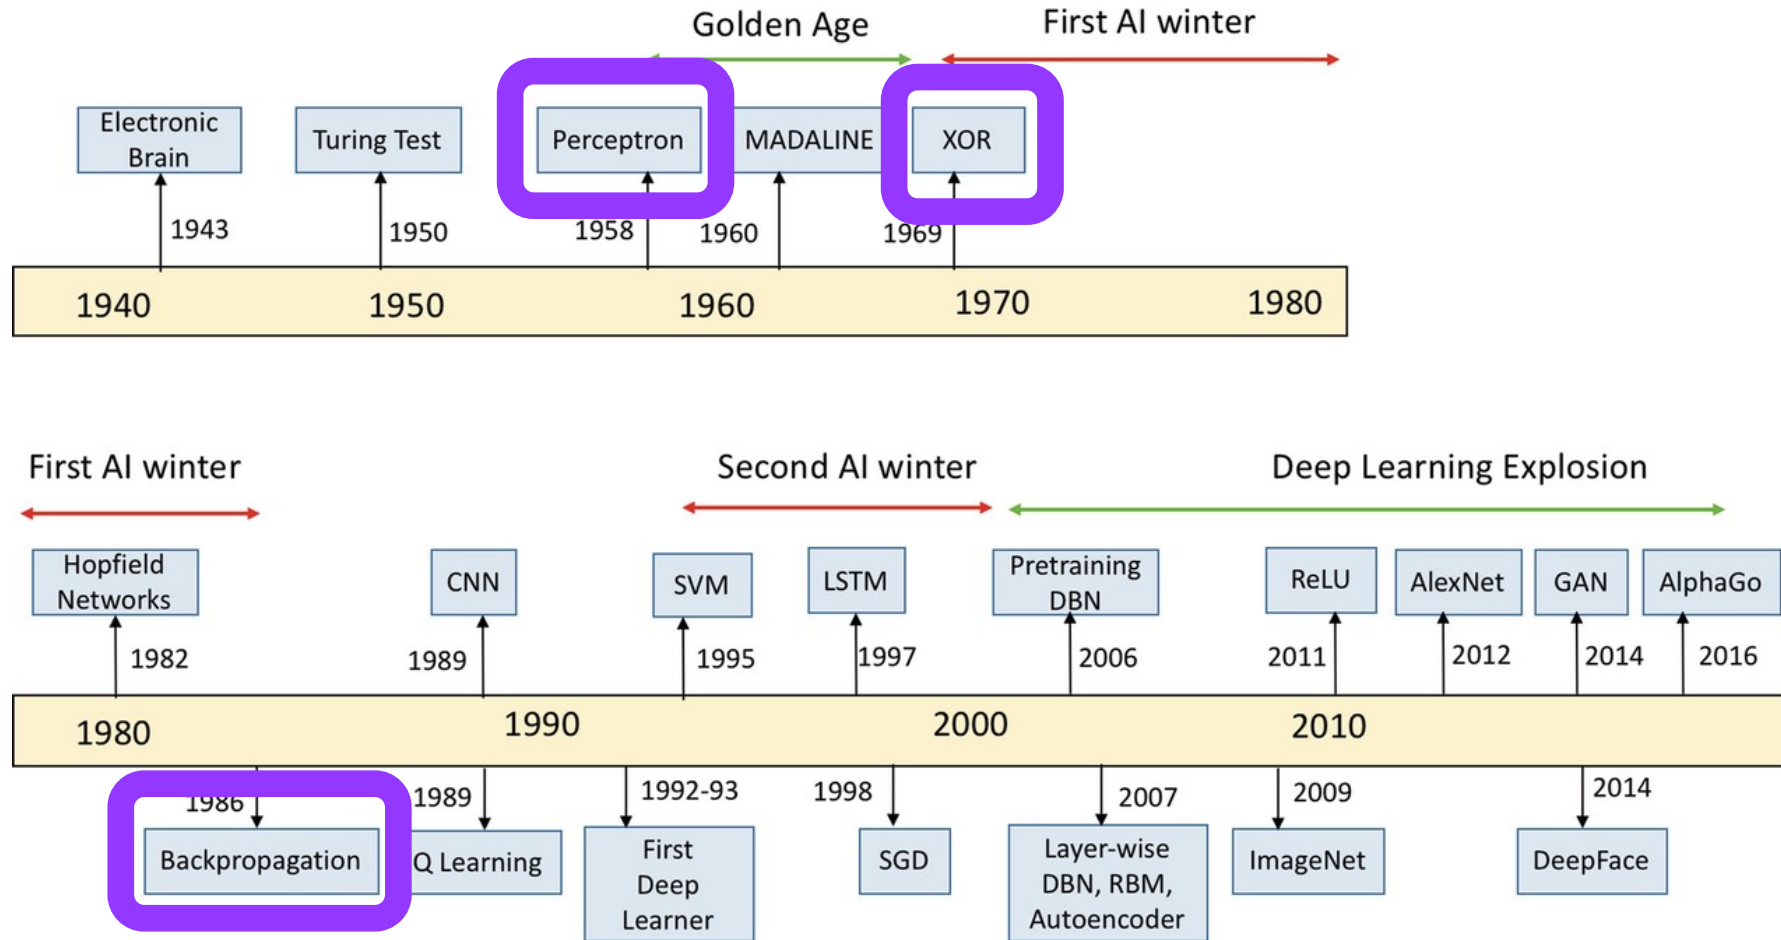
466/566 Fall 2022

Some of these slides based on content from Seyong Kim and Nando de Freitas  
Nando's youtube lectures (which are really great!) See last slide for link

# Administrivia

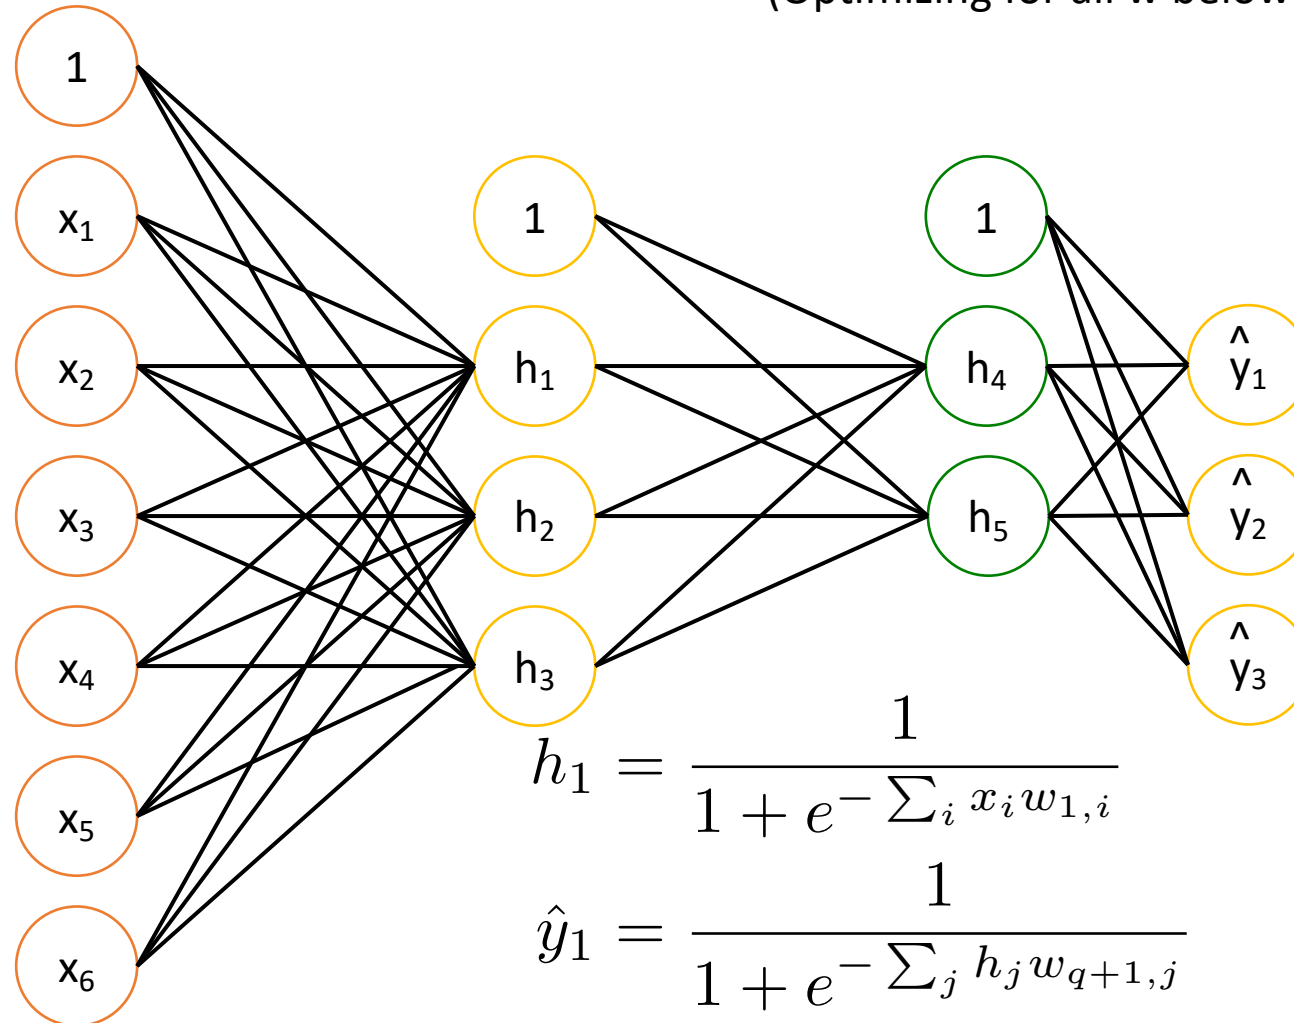
- Monday is a holiday (**no office hours, no lab**)
- Thursday of next week (Oct 13): **no class**
- Midterm is Oct 20 (two weeks from today)
- We will have some time for **review in class on the 18<sup>th</sup>**
  - Come with questions, I will not be preparing anything

A little of the history of NNs



# Now: Back Prop

(Optimizing for all w below)



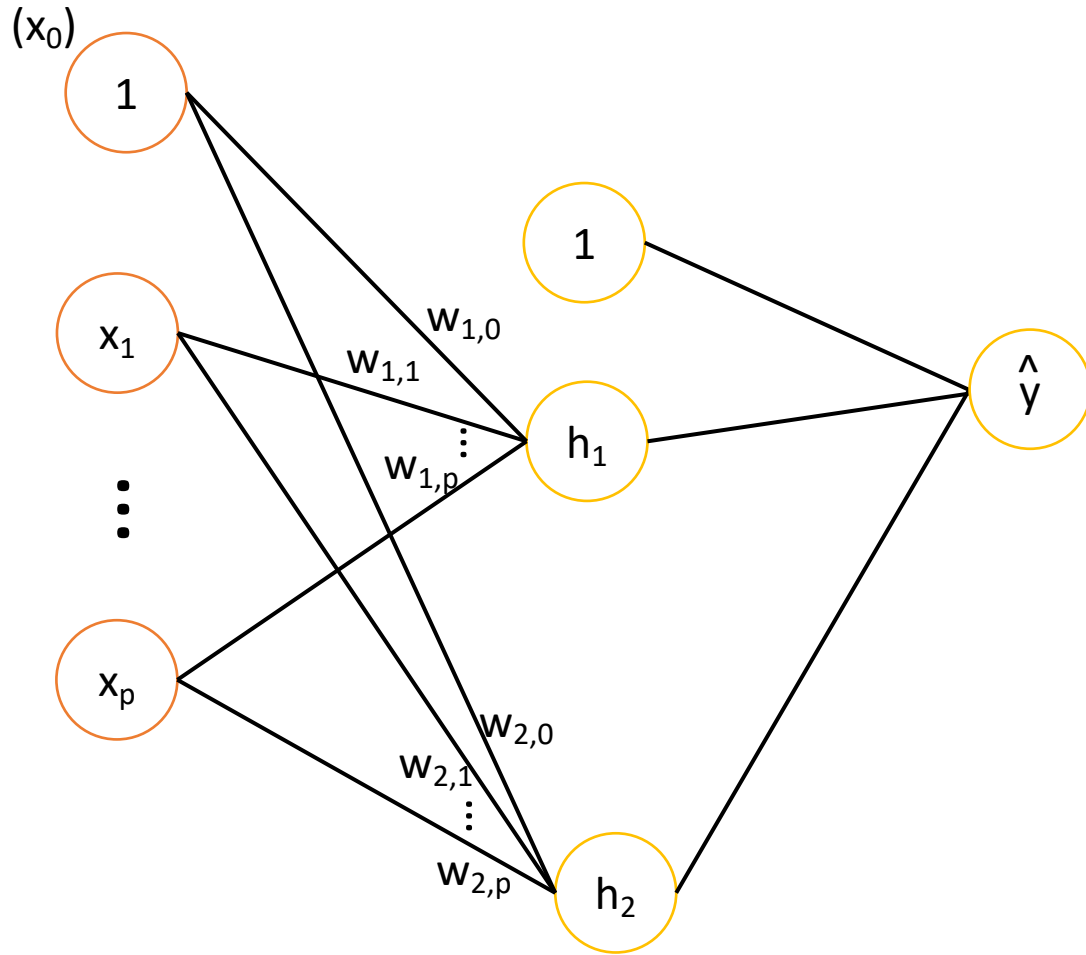
q is total # of hidden nodes across all layers

# Backprop Algorithm

1. Randomly initialize weights ( $w$ )
2. Repeat until convergence
  1. For each (batch, mini-batch) in data
    1. For each data point in batch
      1. Forward pass (calculate all intermediate values  $h$ ,  $s$ )
      2. Backwards pass compute gradient of loss wrt all  $w$
    2. Average gradient over data points in batch
  3. Update  $w$

Backprop takes advantage of the fact that many of the value you need for each gradient are computed in the forward pass, or as part of another gradient.

# Our example for in class



(Optimizing for all w below)

$$h_1 = \frac{1}{1 + e^{-\sum_i x_i w_{1,i}}}$$

$$\hat{y} = \sum_j h_j w_{3,j}$$

$$\hat{y} = \sum_j h_j w_{3,j}$$

$$h_1 = \sigma(s_1)$$

$$s_1 = \sum_{i=0}^p w_{1,i} x_i$$



$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

# Techniques for speeding up SGD





# Momentum

- Sometimes SGD updates produce oscillations, overstepping the most optimal path.
- Momentum mixes a fraction of the last update into the current update, which helps control oscillations



# Momentum

- $v_t = \gamma v_{t-1} + \eta \nabla_w J(w)$
- $w = w - v_t$
- The update is a mix of the regular SGD update ( $\eta \nabla_w J(w)$ ) and the last update ( $\gamma v_{t-1}$ )  $\gamma$  usually around 0.9



# Nesterov accelerated gradient (NAG)

- Momentum update can be rewritten

- $w = w - \gamma v_{t-1} - \eta \nabla_w J(w)$

Lookahead

- NAG notes that we have some of the info we need to compute part of that update ahead of time

- $v_t = \gamma v_{t-1} + \eta \nabla_w J(w - \gamma v_{t-1})$

- $w = w - v_t$

# Adagrad

- Adapt the learning rate based on the frequency of a feature
  - Features that appear often have weights that are updated with a smaller step size
  - The *learning rate* changes based on magnitude of the past updates
- $\mathbf{g}_t = \nabla_{\mathbf{w}} J(\mathbf{w}_{t,i})$
- $\mathbf{w}_{t+1} = \mathbf{w}_t - \frac{\eta}{\sqrt{G_t + \epsilon}} \odot \mathbf{g}_t$ 

Without the denominator, this is just the regular SGD update
- $G_t$  is a diagonal matrix where each diagonal element  $i,i$  is the sum of the squares of the gradients w.r.t. element  $i$  of  $\mathbf{w}$  up to time step  $t$

# Adagrad

- $w_{t+1} = w_t - \frac{\eta}{\sqrt{G_t + \epsilon}} \odot g_t$
- As  $t$  grows (more and more epochs) the scaling of the gradient update can become very aggressive, meaning that updates basically stop!

# Adadelta

- Adagrad scales the learning rate by a factor of *all* past updates
- Solution: scale only by a *window* of past updates
  - Avoid storing all past updates in the window by applying a multiplier  $g < 1$
  - Updates  $t$  time steps away will be decayed by  $g^t$ , driving down the contribution of updates that were far in the past.

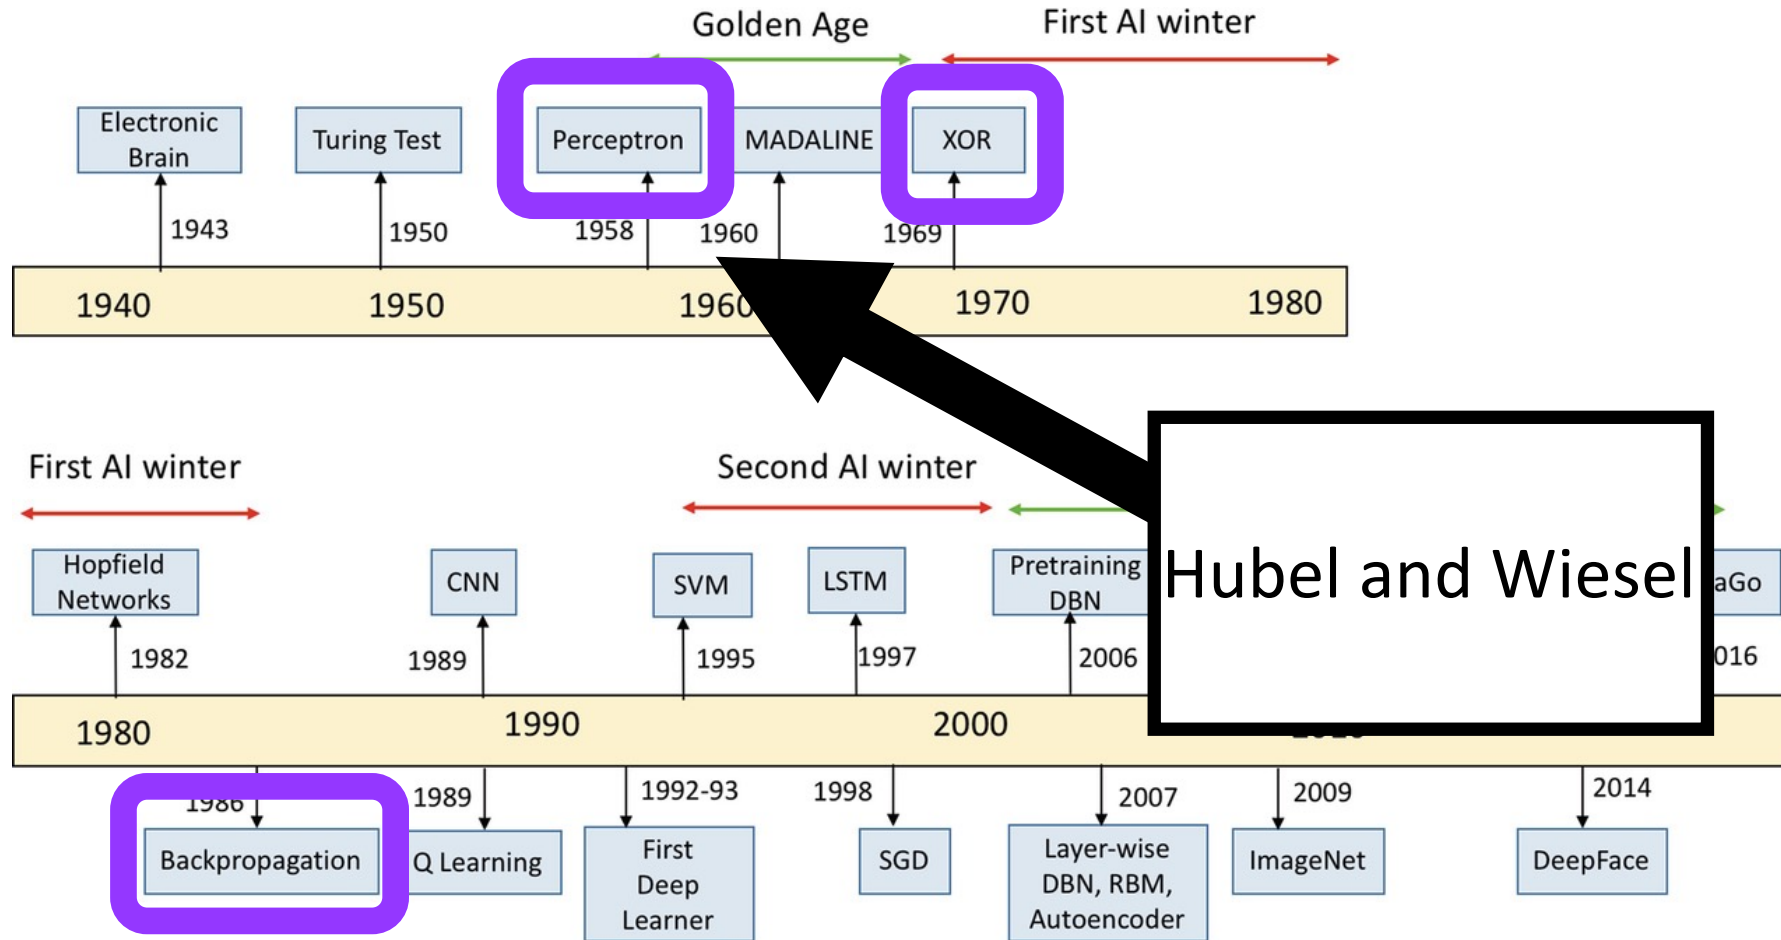
# ADAM

- A mix of momentum and adadelta

# Second-order methods

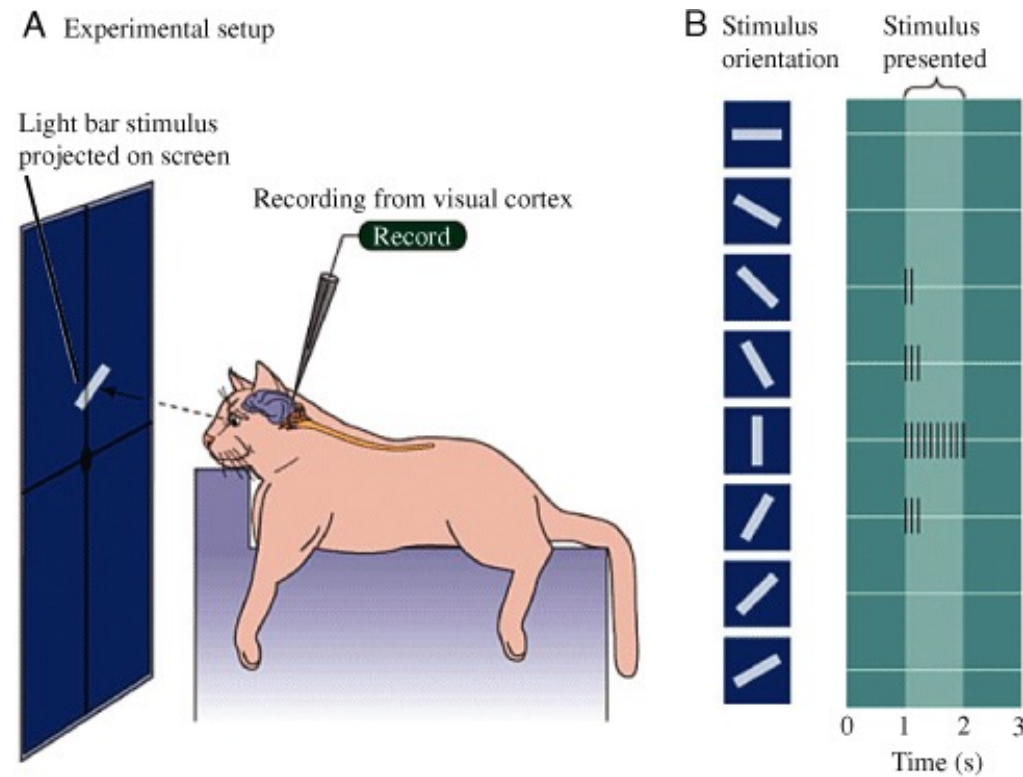
- Not often used because
  - Requires computation and storage of all 2<sup>nd</sup> order derivatives (M params ->  $M^2$  2<sup>nd</sup> order derivatives)
  - Approximated with full dataset (which are typically v. large in DL)
- Some methods to approximate this with less overhead
  - L-BFGS

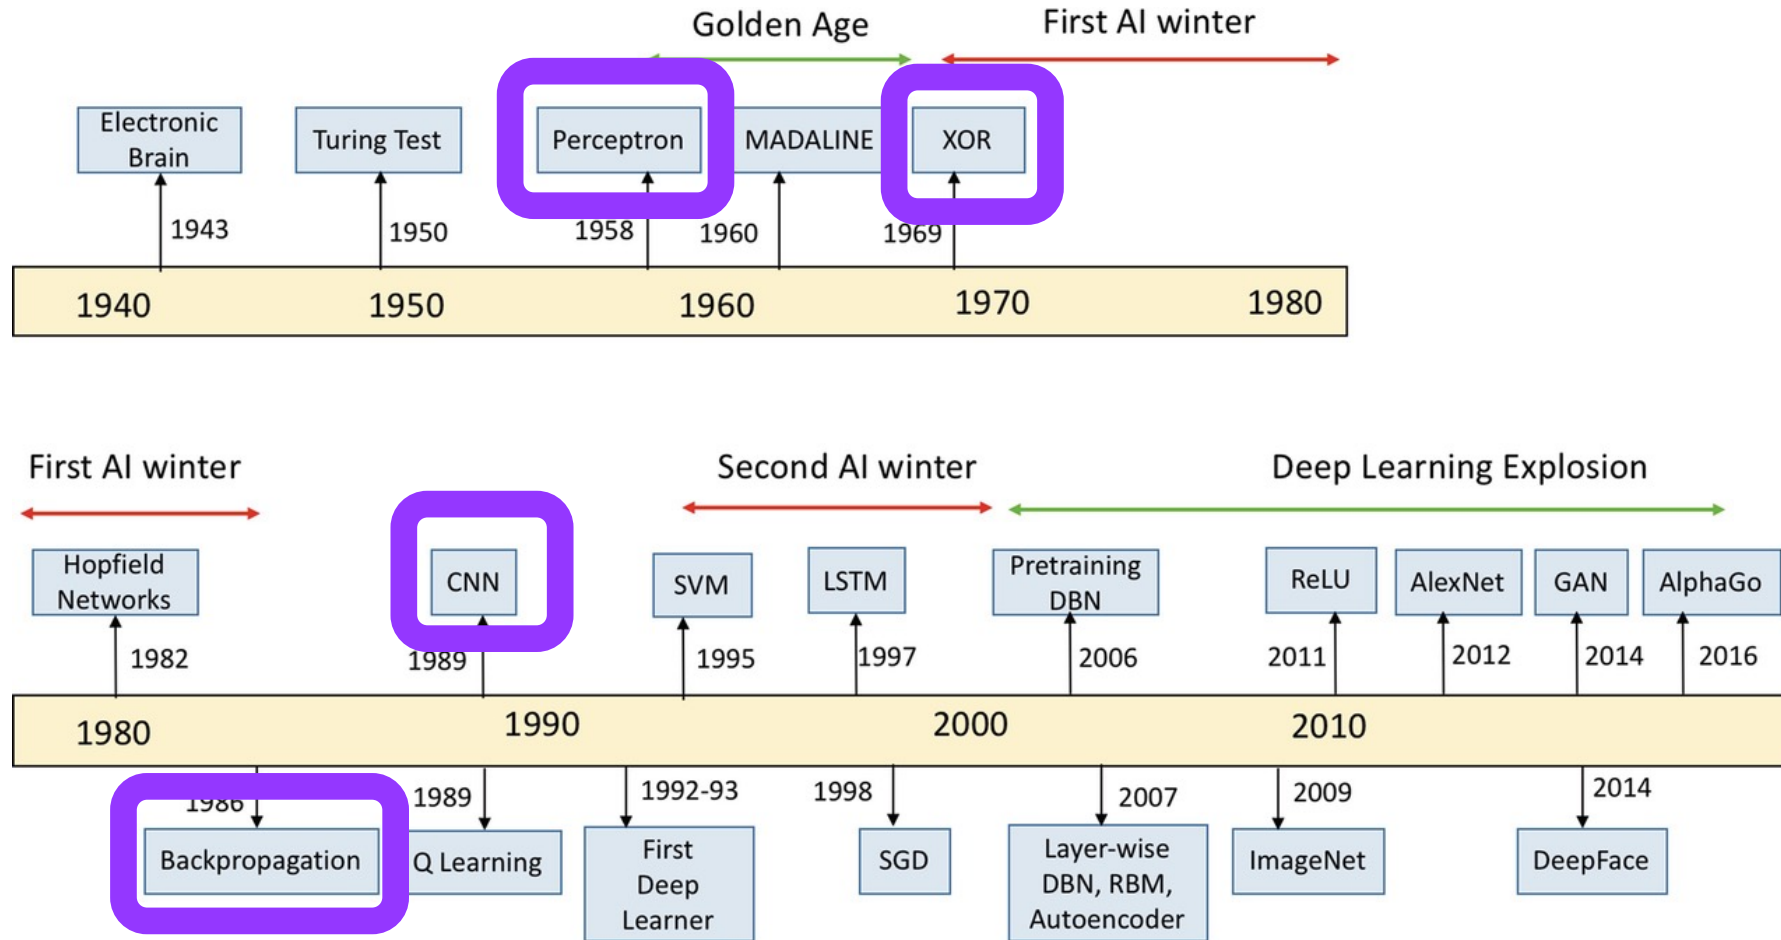




# An Important Discovery

- In this video, the static noise you hear is a representation of the neurons firing in response to the visual stimulus
- <https://www.youtube.com/watch?v=jw6nBW021Zk>





# CNNs: Convolution

- Hubel and Wiesel inspire the idea of convolution in neural networks
  - The same edge detector behavior can be found in multiple receptive fields
- CNNs are powerful because the same “**filter**” (i.e. edge detector) is repeatedly used on all patches of the image
  - This saves parameters, makes learning more efficient


# CNNs: Convolution

- Output of convolution




# What do CNNs learn?

- When trained on images



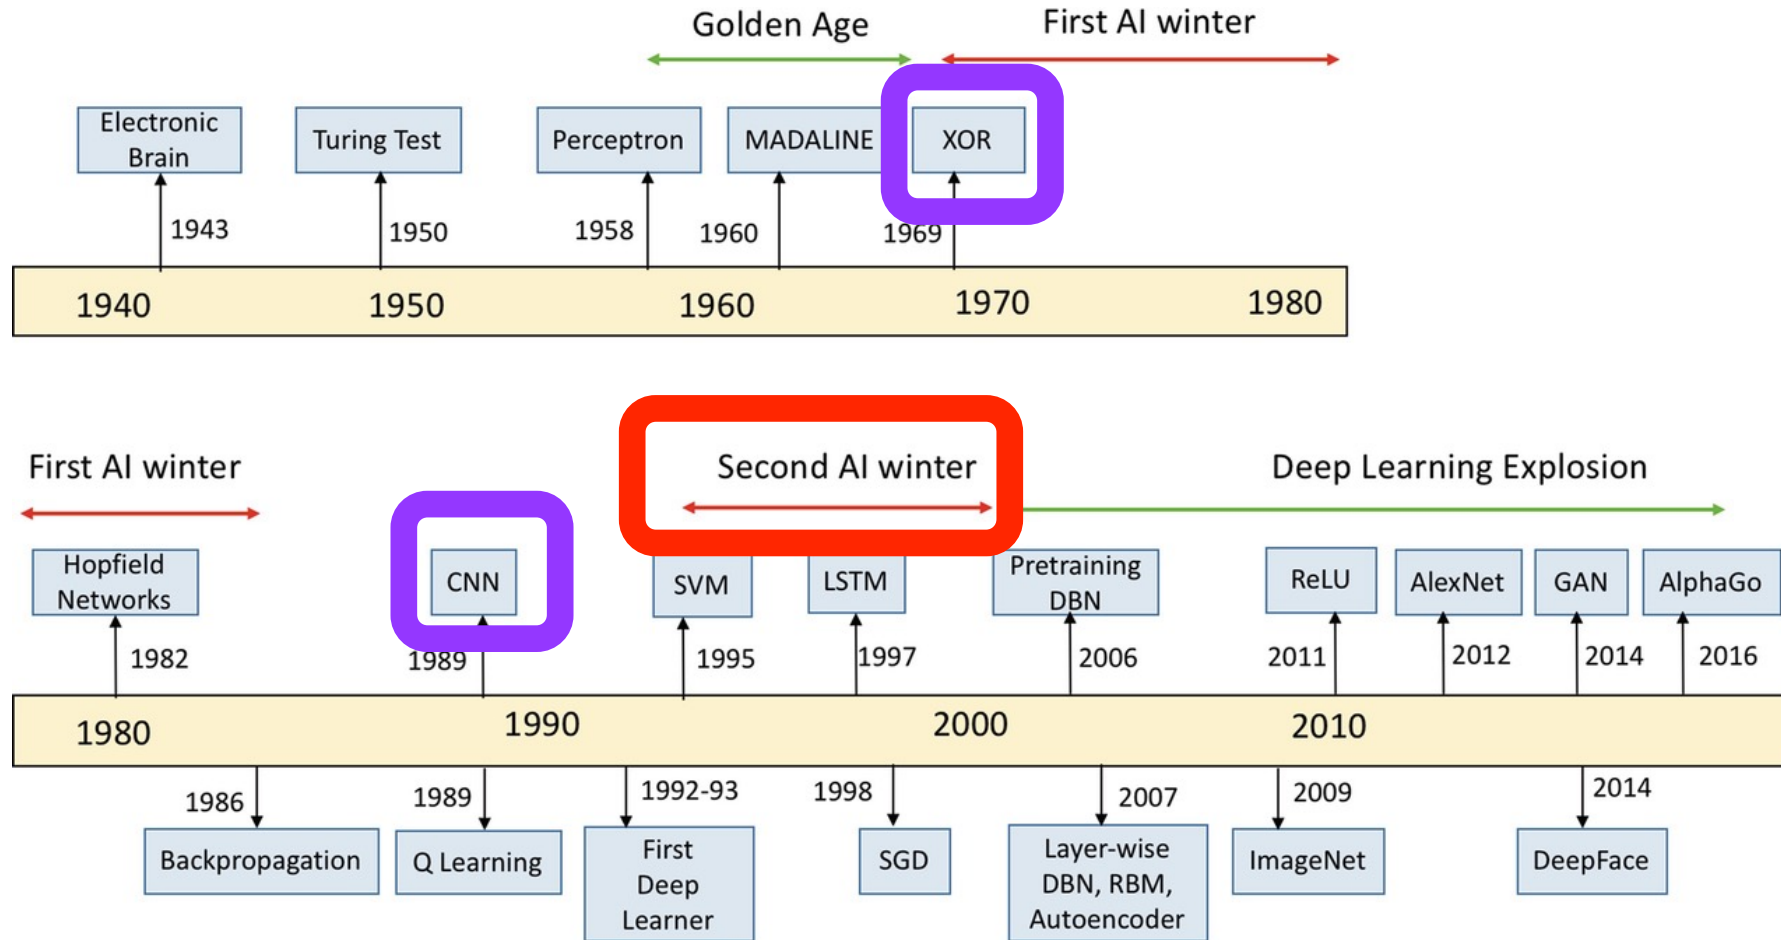
# What do CNNs Learn?



# CNNs

- CNNs were extremely useful for simple tasks
  - E.g. character recognition for hand written digits
- But, CNNs couldn't handle more complex problems
  - There wasn't enough data
  - Computers weren't powerful enough





# The Second AI Winter

- AI Hype grew and grew
- Expert systems became very popular
  - Used databases of knowledge to mimic human decision making
- Companies stated “We’ve built a better brain” and declared that “[I]t is now possible to program human knowledge and experience into a computer ... Artificial intelligence has finally come of age.”

# The Second AI Winter

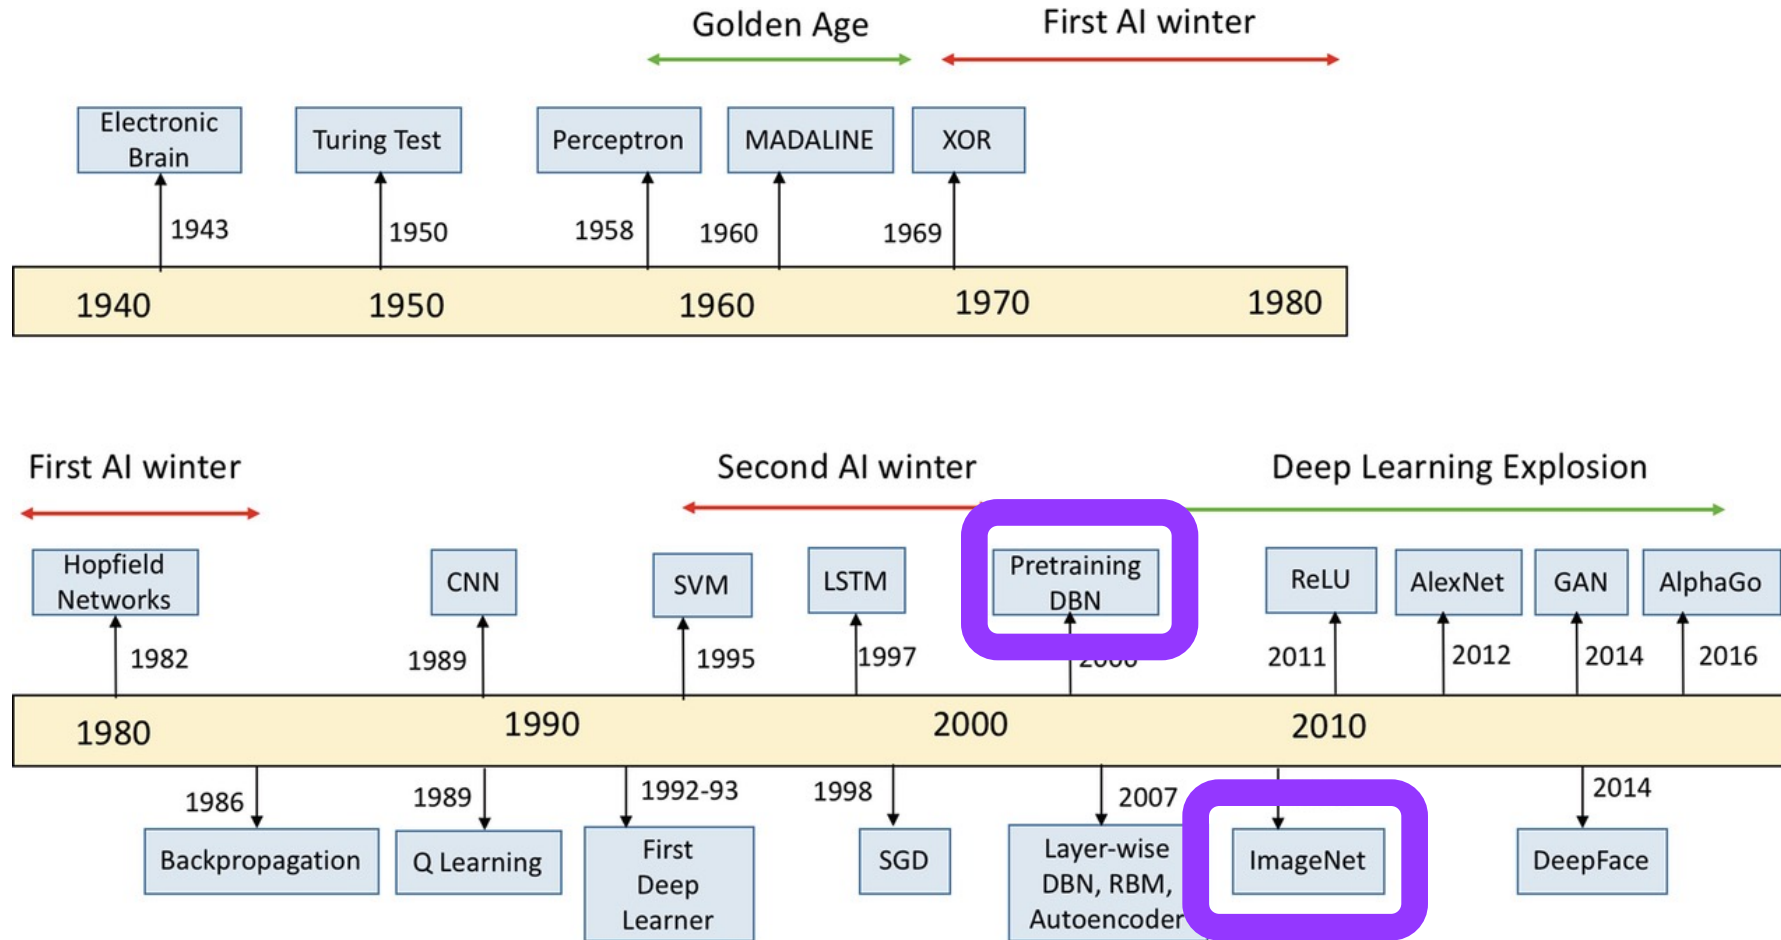
- People became skeptical
- [John McCarthy] described the expert system MYCIN built to assist physicians.
  - He then laid out a situation where a patient has *Cholerae Vibrio* in his intestines.
  - When asked, the systems prescribed two weeks of tetracycline.
  - This would most likely kill off all the bacteria, but by then the patient would already be *dead*.

# The Second AI Winter

- The databases of “human knowledge” in expert systems had to be created manually
  - Rules to operate over these databases also had to be manually made
- Many tasks are too complicated for engineers to **design rules** for manually.
  - E.g. Systems for vision, medical diagnostics, etc


# The Second AI Winter

- The general interest in AI declined as the expectations could not be met.
- Many AI companies closed their doors.
- The AAAI conference that attracted over 6000 visitors in 1986 quickly decreased to just 2000 by 1991.





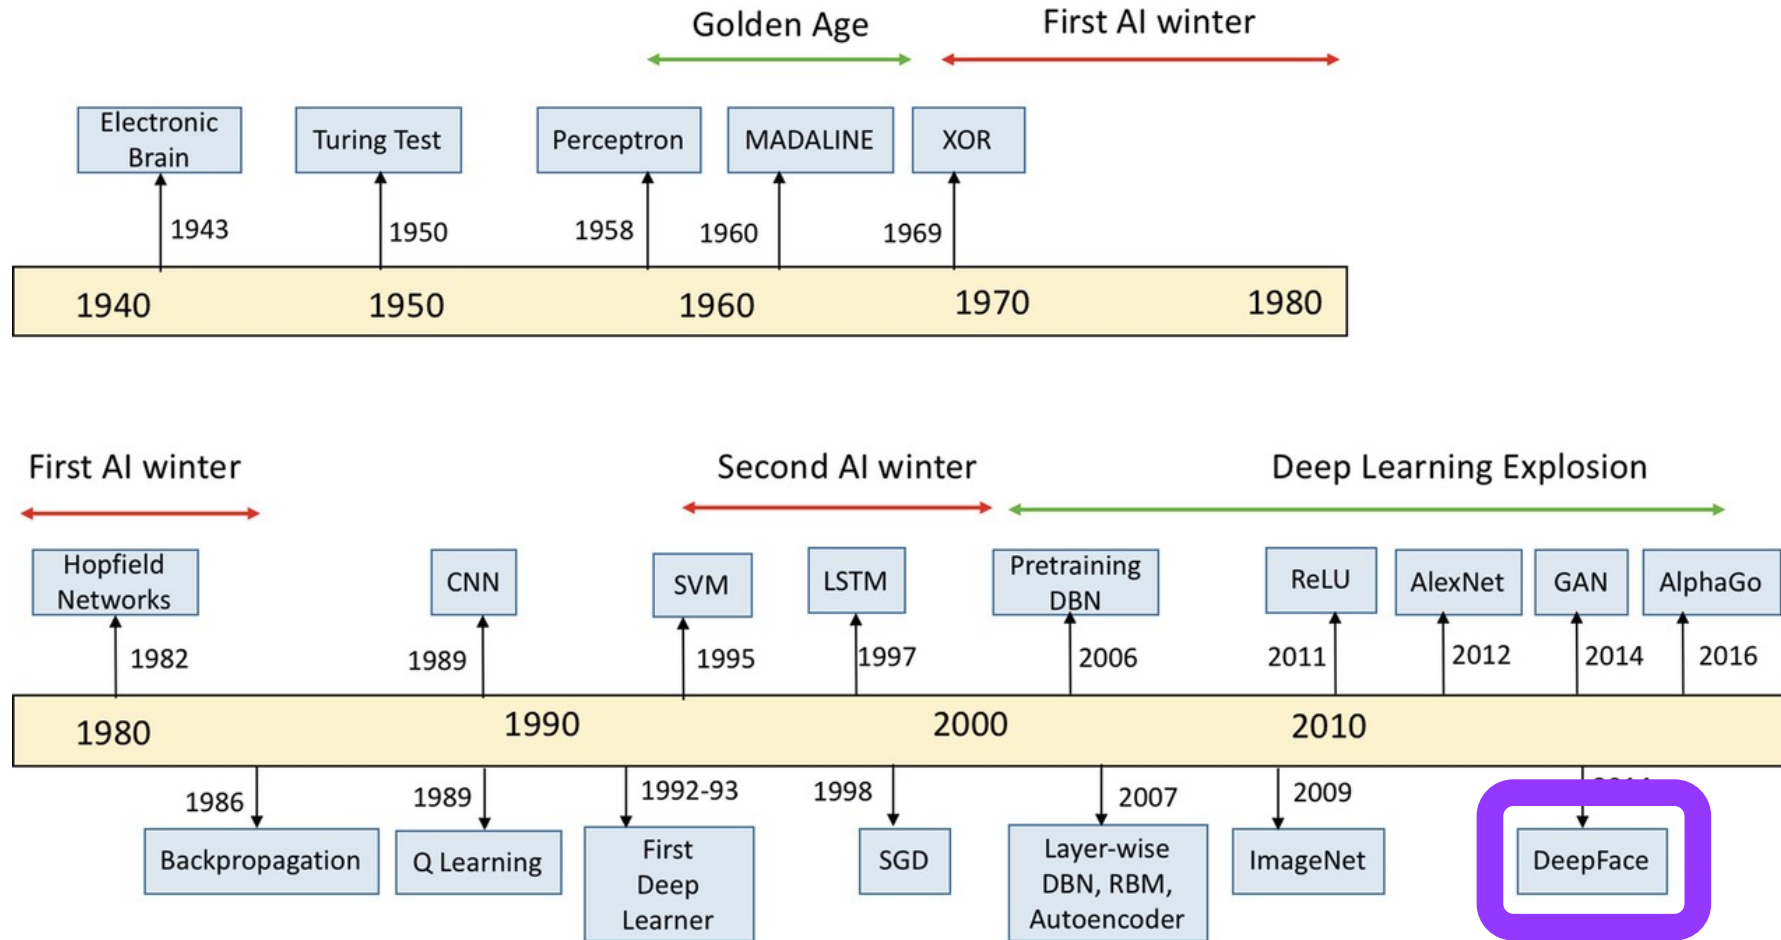
# ImageNet



# ImageNet Enables Learning

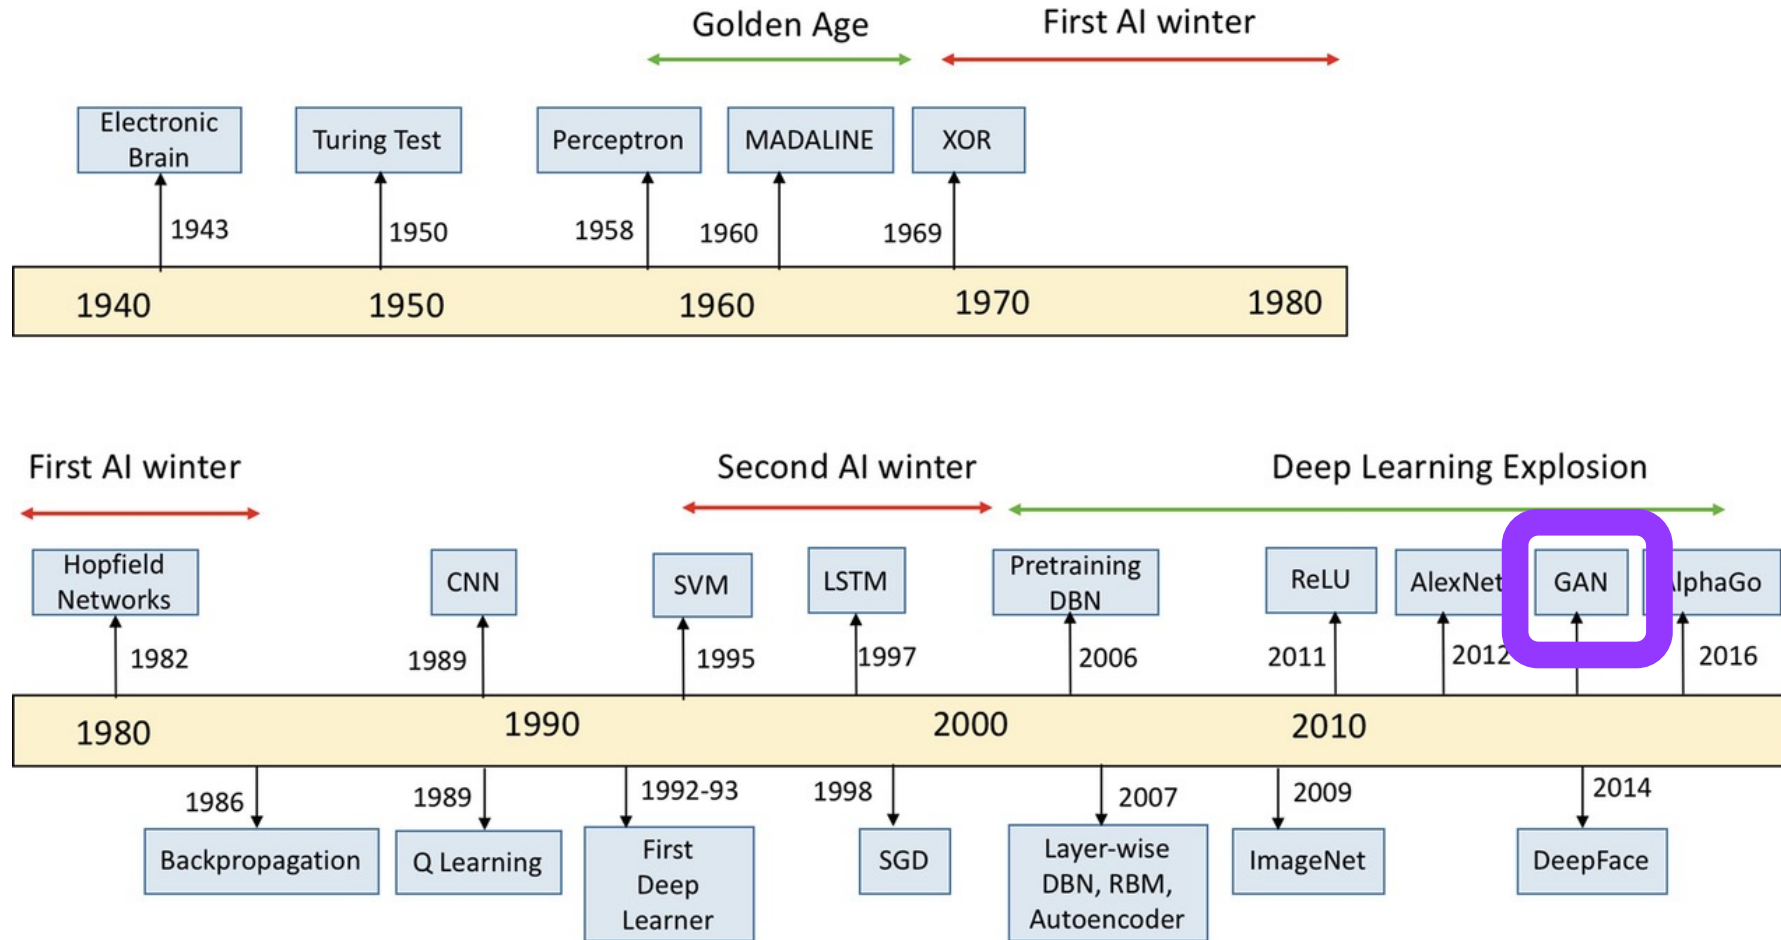
- The size of ImageNet, and the increasing speed of computers, allows for CNNs to become world class object detectors!
- In contrast to expert systems, CNNs learn their database of filters, and the functions (rules) that operate over them
  - Much more powerful
  - Generalize well to novel images
  - Generalize well to new problem domains





# DeepFace




- Facebook used images **uploaded and tagged by its users** to build a face recognition system with 97.3% accuracy
- At the time, was the largest facial dataset to-date, an identity labeled dataset of four million facial images belonging to more than 4,000 identities.



# First: Adversarial Examples

# Adversarial Examples



- We can make small changes to an image (nearly imperceptible!) and cause a network to misclassify
- Extreme implications for e.g. self driving cars



correct

+distort



ostrich



correct

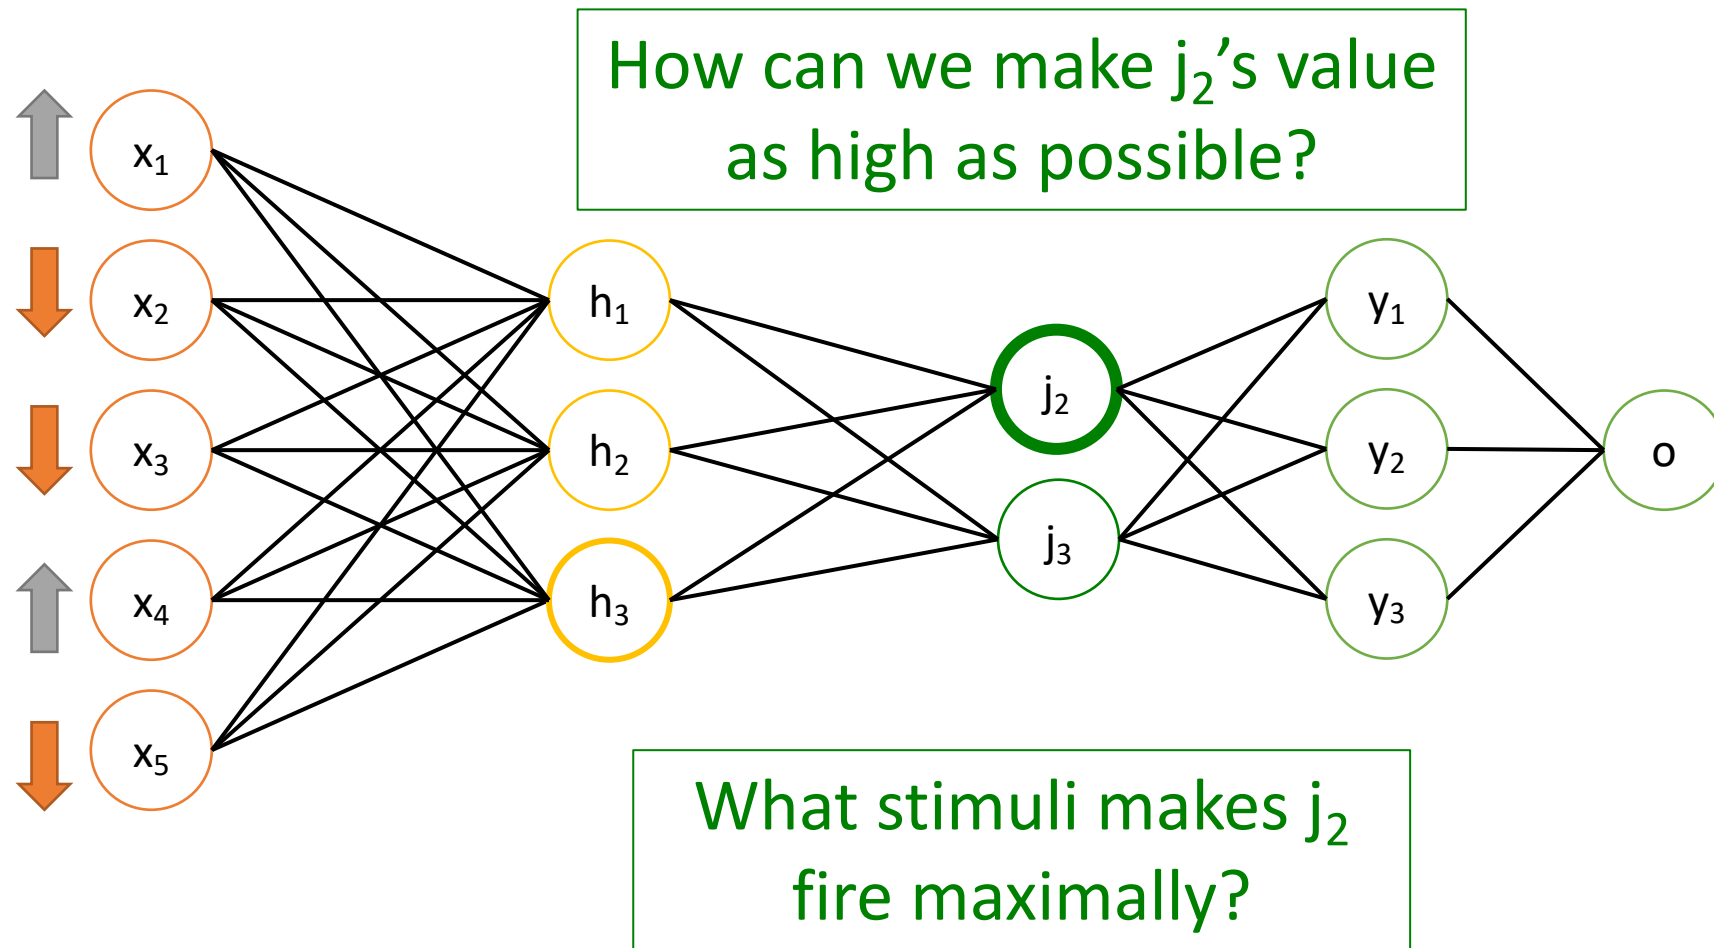
+distort

ostrich




How are adversarial examples made?

Recall:






CNN3



CNN2




CNN1




CNN8



CNN7



CNN6

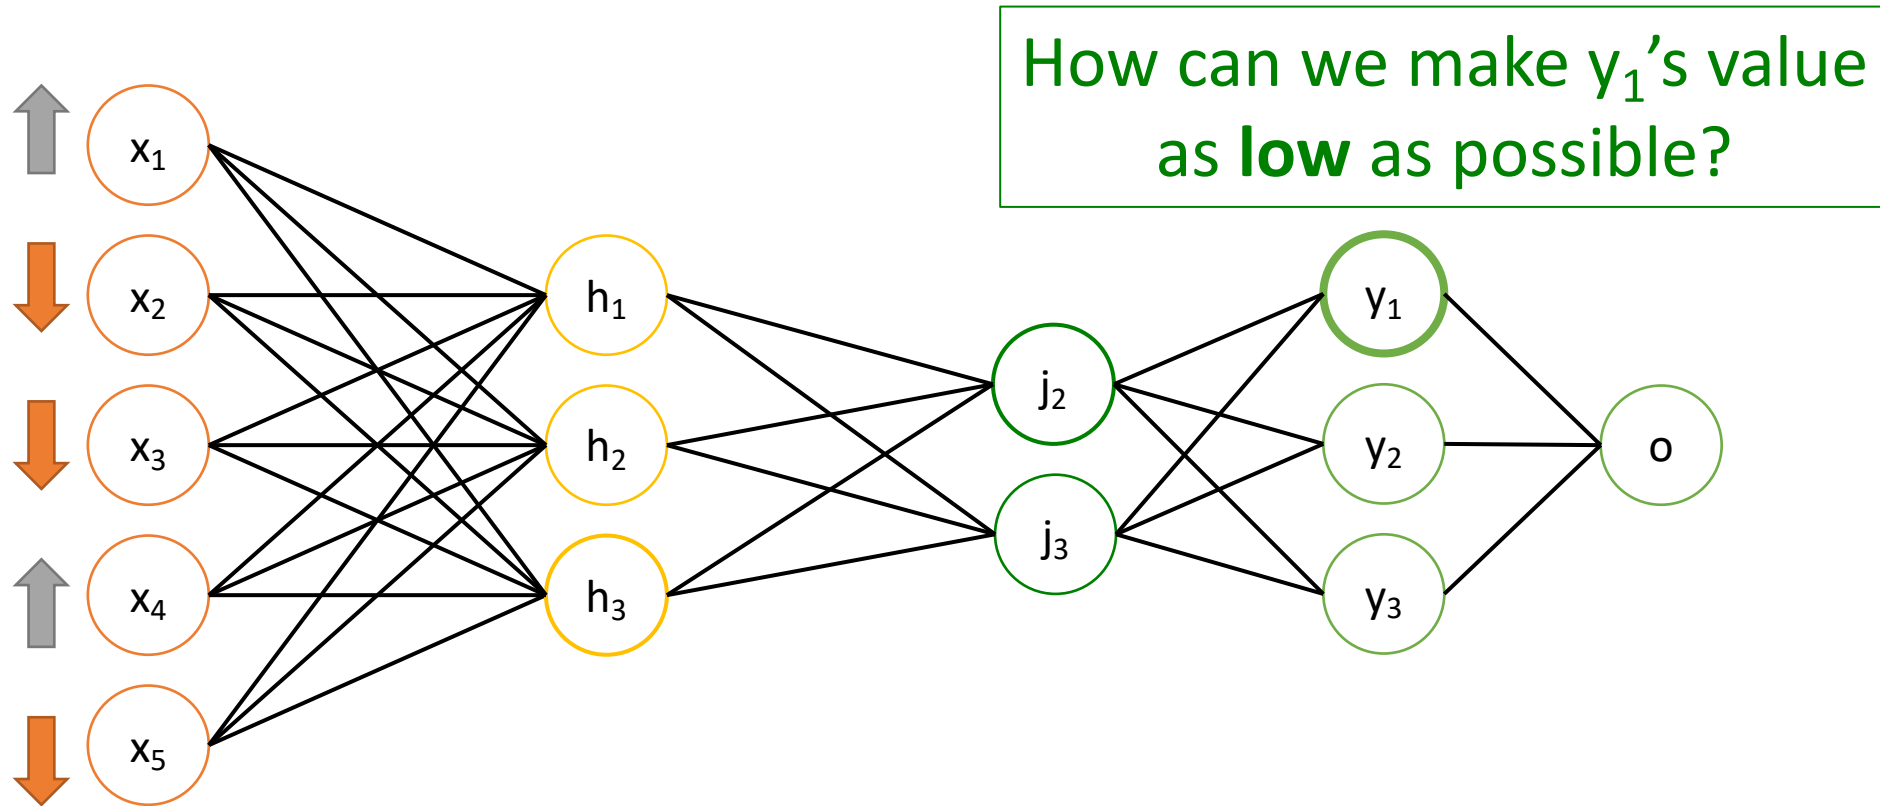


Horikawa & Kamitani (2017)

# How are adversarial examples made?

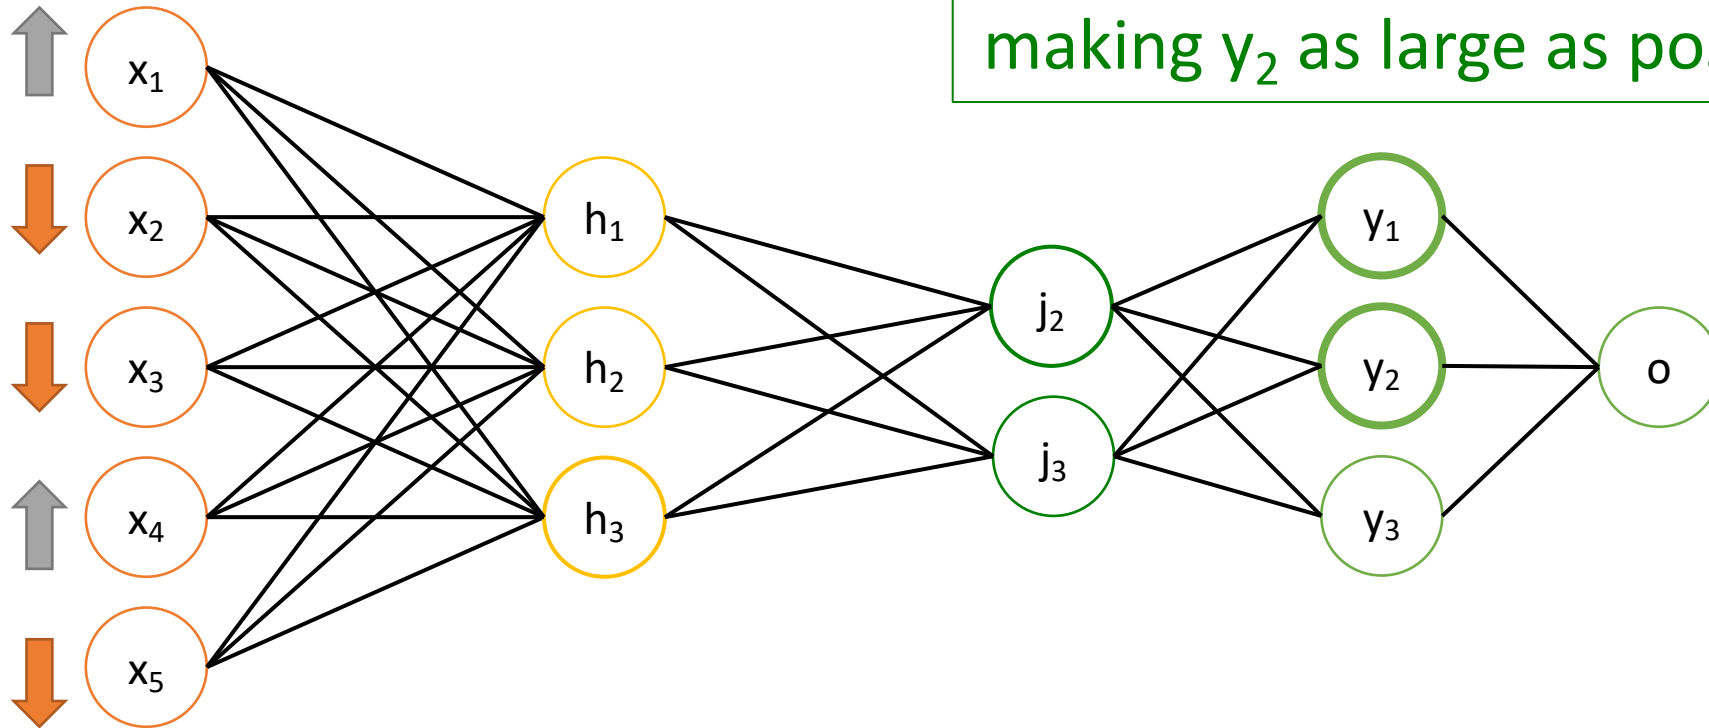
- An adversarial example forces the model to predict the wrong class
  - Sometimes a specific wrong class

Correct class is  $y_1$



Correct class is  $y_1$

How can we make  $y_1$ 's value as **low** as possible while also making  $y_2$  as large as possible?



# Removing Stop Signs with Stickers



# Making Stop Signs with Stickers





# Making Stop Signs with Stickers



# GAN

- Generative Adversarial Network
- Two dueling neural networks
  - One trained to generate images
  - One trained to distinguish generated images from true images



# GAN




As training progresses, the generator gets closer to producing output that can fool the discriminator:



Finally, if generator training goes well, the discriminator gets worse at telling the difference between real and fake. It starts to classify fake data as real, and its accuracy decreases.



# GAN







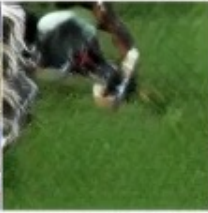















# GAN



Figure 1: Class-conditional samples generated by our model.

# GAN

Real Image	P-AttnGAN	P-AttnGAN w/ Lyt	Obj-GAN w/ SN	Obj-GAN
				
A brown dog lying on bed with his banana toy.				
				
A black and white dog catching a frisbee on a field.				
				
A couple of elephants standing by some trees.				
				
Several sheep dotting a grass hillside near a mountain edge.				



# GAN



# DeepFace or DeepFake?

- <https://www.creativebloq.com/features/deepfake-examples>

# Resources

- Programming resources for training your own NNs
  - Tensorflow <https://www.tensorflow.org/>
  - Keras <https://keras.io/>
  - Pytorch <https://pytorch.org/>
  - For intuition: <http://playground.tensorflow.org/>
- Short course on deep learning (Nando De Freitas)
  - <https://www.youtube.com/playlist?list=PLjK8ddCbDMphlMSXn-w1IjyYpHU3DaUYw>
- Commentary on AlphaGo
  - <https://www.youtube.com/watch?v=UMm0XaCFTJQ>
  - <https://www.youtube.com/watch?v=g-dKXOlsf98>
- Other fun videos
  - Geoff Hinton is in this one! Neural Net stuff is towards the end
    - <https://www.youtube.com/watch?v=yxxRAHVtafl>
  - Fei Fei Li's Ted Talk (Creator of ImageNet)
    - [https://www.ted.com/talks/fei\\_fei\\_li\\_how\\_we\\_re\\_teaching\\_computers\\_to\\_understand\\_pictures?language=en](https://www.ted.com/talks/fei_fei_li_how_we_re_teaching_computers_to_understand_pictures?language=en)